

Fish Detection

Yolov3 in an Industrial Env

Summary

Project and Data Introduction

Implementation pipeline

YoloV3 Architecture

Training Results

YoloV3 Inference Results

Project and Data Introduction


In this project, we are working on detecting bigeye and yellowfin tunas on video feeds from fishing vessels. The task shall also include classifying the fish, as well as provide a rectangular bounding box around each fish.

```
img_id,bbox_id,x_min,x_max,y_min,y_max,label_name
94a95d88-23f0-11e9-ae08-6377a4c78e35,1,392,607,197,405,Human
94a95d88-23f0-11e9-ae08-6377a4c78e35,2,315,444,132,260,Human
94a95d88-23f0-11e9-ae08-6377a4c78e35,3,228,353,36,183,Human
94a95d88-23f0-11e9-ae08-6377a4c78e35,4,166,331,286,470,Human
94a95d88-23f0-11e9-ae08-6377a4c78e35,5,331,377,359,513,Yellowfin tuna
94a95d88-23f0-11e9-ae08-6377a4c78e35,6,156,191,513,589,Yellowfin tuna
94a95d88-23f0-11e9-ae08-6377a4c78e35,7,531,605,0,71,Human
94a95d88-23f0-11e9-ae08-6377a4c78e35,8,50,150,95,217,Human
94a95d88-23f0-11e9-ae08-6377a4c78e35,9,0,51,226,306,Human
94ac6e38-23f0-11e9-b635-cb6efd3bfc7b,1,318,625,367,488,Black marlin
94ac6e38-23f0-11e9-b635-cb6efd3bfc7b,2,120,291,361,472,Human
04a96a28-23f0-11e9-b635-cb6efd3bfc7b,3,250,132,238,132,Human
```

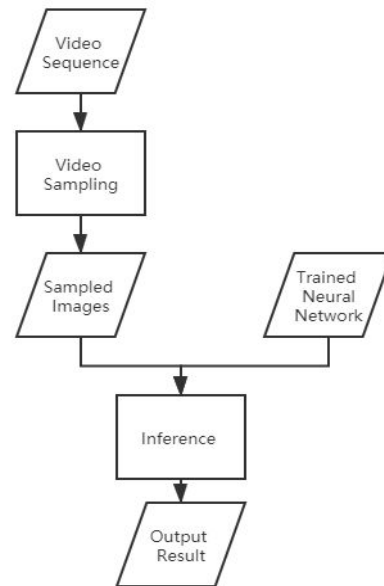


Implementation Pipeline

Training Phase

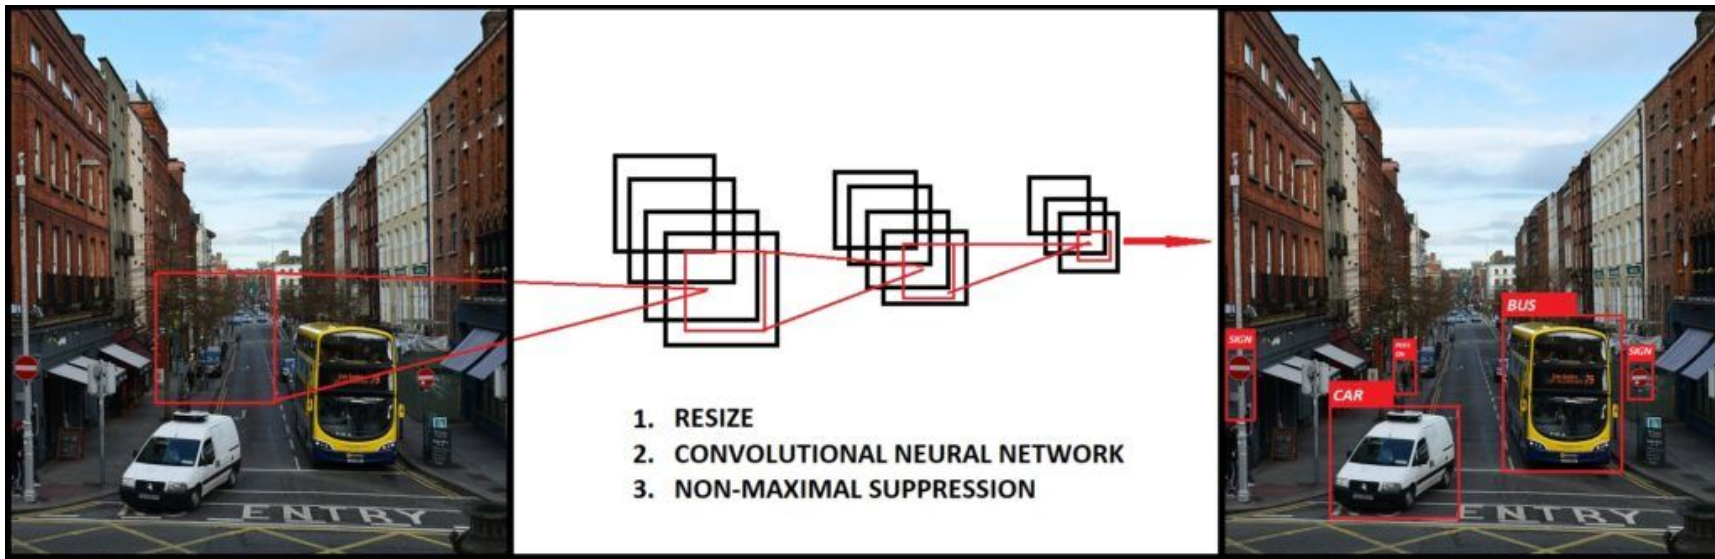


Fine Tuning Phase



YouOnlyLookOnce V3

Yolo is a state-of-the-art object detection system (network). It was developed by Joseph Redmon and it is known for its speed, being much faster than R-CNN and others, we can achieve real-time object detection.




YouOnlyLookOnce V3

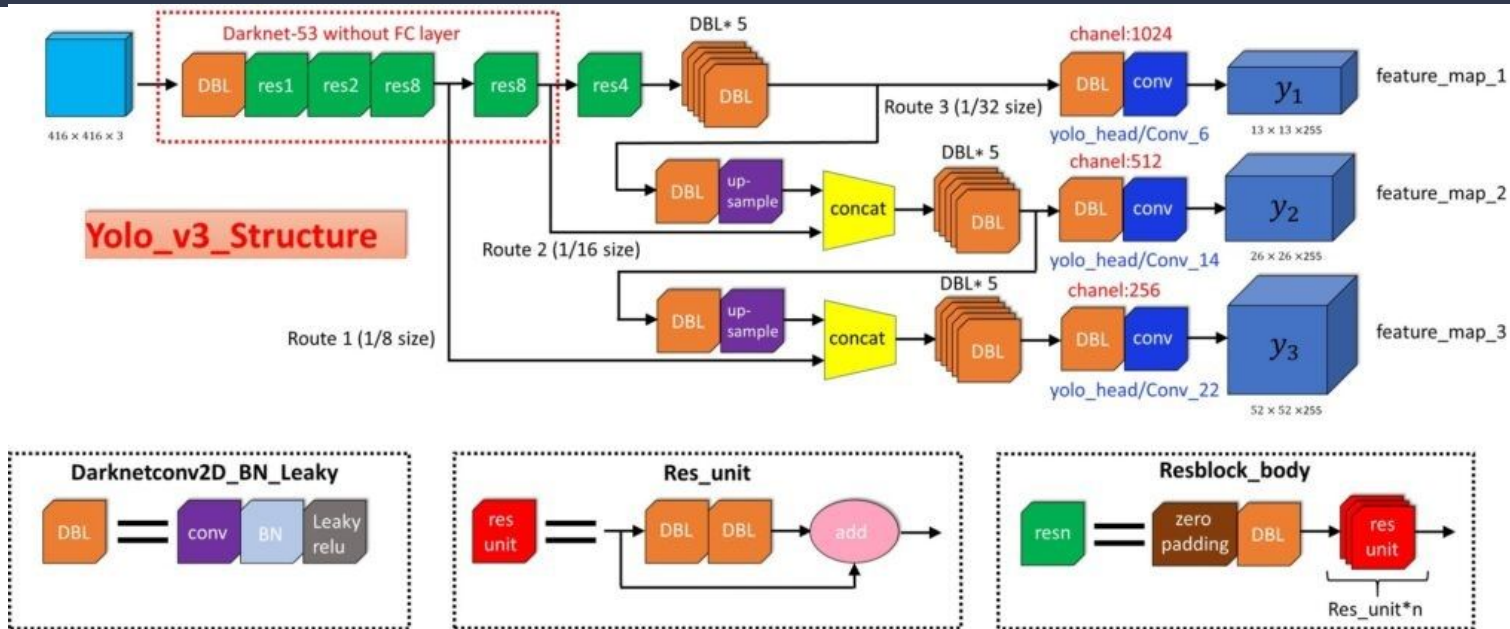
YOLO has reframed an object detection problem into a single regression problem.

It goes directly from image pixels, up to bounding box coordinates and class probabilities.

Hence, a single convolutional network predicts multiple bounding boxes and class probabilities for those boxes.



YouOnlyLookOnce V3 Architecture



Output Format : $S \times S \times [B * (4 + 1 + N^{\circ} \text{Classes})]$


S : grid size (7 in our case)

B : N° Scales (YoloV3 uses 3)


4 : Predicted Bounding Box Coords

1 : Objectness of predicted BB

Training Results



Inference Results



Thank You,

Any Questions ?

A dark blue diagonal gradient bar that starts from the bottom left and extends towards the top right, covering the lower half of the slide.